

Industry-Scale Knowledge Graphs: Lessons and Challenges

Julia Grace Muddada
Knowledge Engineering
College of Engineering and Computer Science

KEY TAKEAWAYS - 1

Knowledge Graphs are Central but Domain-Specific

Each company designs its KG around its core needs—

Microsoft for search and professional networks,

Google for universal search,

Facebook for social context,

eBay for product catalogs,

IBM for enterprise discovery.

Despite different focuses, all treat KGs as foundational to powering intelligent applications.

KEY TAKEAWAYS - 2

Building and Maintaining KGs is Technically Challenging

Common difficulties include Entity disambiguation, Knowledge evolution, Knowledge extraction and Scale & Performance.

Companies adopt different architectures (graph stores, hybrid systems, replicated logs) to balance performance, flexibility, and scalability.

KEY TAKEAWAYS - 3

Coverage, Correctness & Freshness

Effective KGs must be broad (coverage), accurate (correctness), and up-to-date (freshness).

Provenance

Provenance tracking is key for user trust, while continuous updates and flexible schemas ensure KGs can evolve with real-world changes.

S U M M A R Y

The paper Industry-Scale Knowledge Graphs highlights how major tech companies like Microsoft, Google, Facebook, eBay, and IBM use knowledge graphs as the foundation for search, recommendations, social platforms, and enterprise discovery.

While each company tailors its KG to specific needs, they all face similar challenges, including entity disambiguation, evolving schemas, handling noisy data, and scaling to billions of entities.

The authors emphasize that the success of a knowledge graph depends on its quality—defined by coverage, correctness, and freshness—as well as the ability to evolve with real-world changes and maintain user trust through provenance and reliability.